

Structured Representations for Knowledge Transfer in Reinforcement Learning

Benjamin Rosman

Mobile Intelligent Autonomous Systems
Council for Scientific and Industrial Research

&

School of Computer Science and Applied Maths
University of the Witwatersrand
South Africa






Robots solving
complex tasks

Large high-
dimensional
action and state
spaces

Many different
task instances

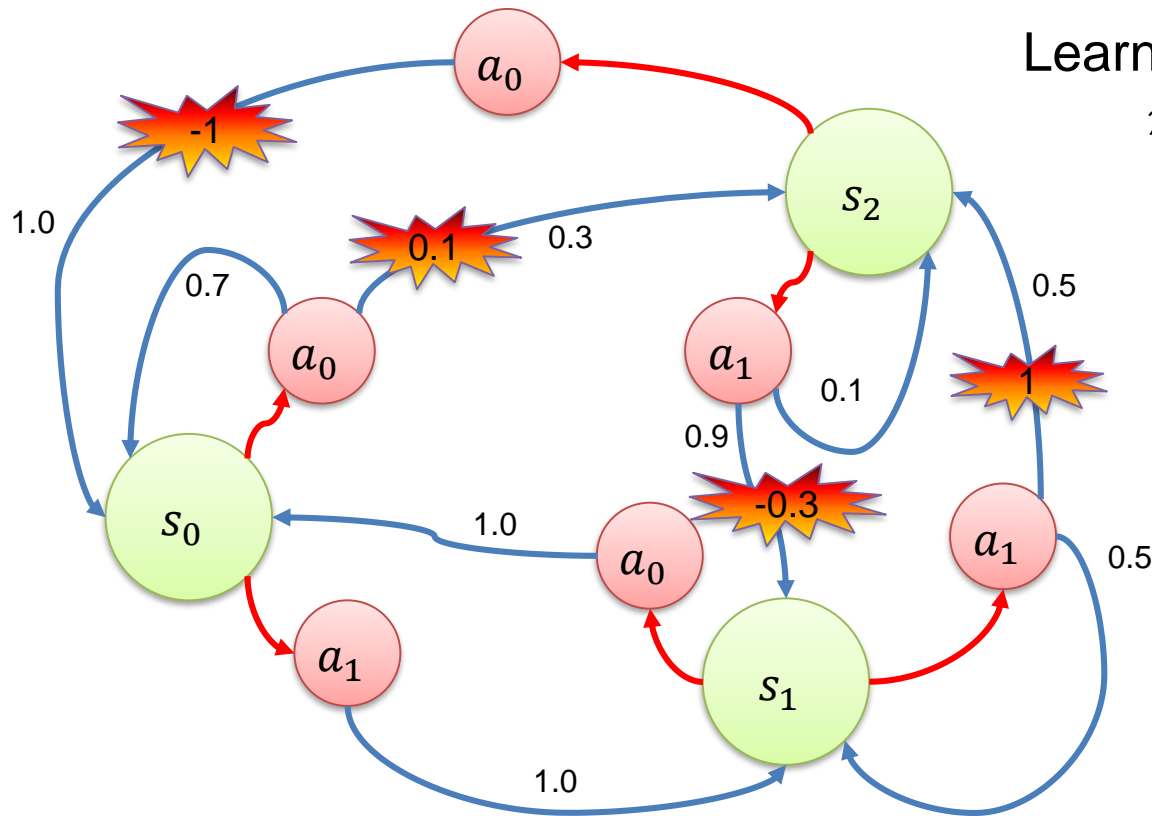
Behaviour learning

- Reinforcement learning (RL)



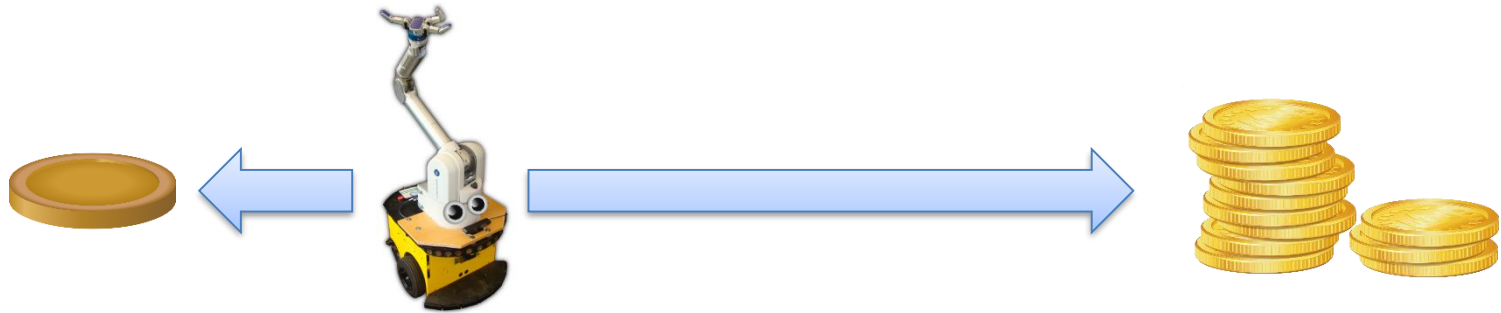
Markov decision process (MDP)

- $M = \langle S, A, T, R \rangle$




Looking into the future

- Can't just rely on immediate rewards




- Define **value functions**:

- $V^\pi(s) = E_\pi\{R_t | s_t = s\}$




- $Q^\pi(s, a) = E_\pi\{R_t | s_t = s, a_t = a\}$



- V^* (Q^*) is a proxy for π^*

Value functions example

- Random policy:




(a)

- Optimal:

3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

(b)



a) gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

b) v_*

→	↕	←	↕	←
↙	↑	↘	←	←
↙	↑	↘	↘	↘
↙	↑	↘	↘	↘
↙	↑	↘	↘	↘

c) π_*



RL algorithms

- So: solve a large system of nonlinear value function equations (Bellman equations)


$$\begin{aligned} V^{\pi'}(s) &= \max_a E \left\{ r_{t+1} + \gamma V^{\pi'}(s_{t+1}) \mid s_t = s, a_t = a \right\} \\ &= \max_a \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma V^{\pi'}(s') \right]. \end{aligned}$$

- Optimal control problem
- But: transitions P & rewards R aren't known!
- RL learning is **trial-and-error learning** to find an **optimal policy** from experience
- Exploration vs exploitation

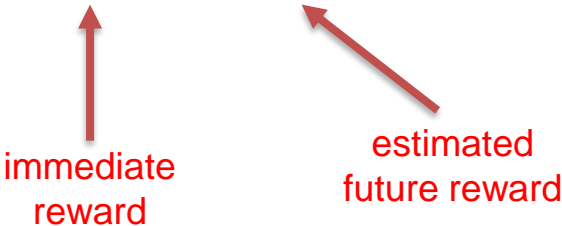
Exploring



Learned value function




An algorithm: Q-learning

- Initialise $Q(s, a)$ arbitrarily
 - Repeat (for each episode):
 - Initialise s
 - Repeat (for each step of episode):
 1. Choose a from s (ϵ -greedy policy from Q)
 - $a \leftarrow \begin{cases} \arg \max_a Q(s, a) & \text{w.p. } 1 - \epsilon \text{ exploit} \\ \text{random} & \text{w.p. } \epsilon \text{ explore} \end{cases}$
 2. Take action a , observe r, s'
 3. Update estimate of Q
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$ learn
 - $s \leftarrow s'$
 - Until s is terminal
- 
- The diagram illustrates the components of the Q-learning update equation. A red arrow points from the text 'immediate reward' to the r term in the equation. Another red arrow points from the text 'estimated future reward' to the $\gamma \max_{a'} Q(s', a')$ term in the equation.

Solving tasks




Generalising solutions?




- How does this help us solve other problems?

Hierarchical RL

- Sub-behaviours: options $o = \langle I_o, \pi_o, \beta_o \rangle$
 - Policy + initiation and termination conditions



- Abstract away low level actions
- Does not affect the state space



Abstracting states

- **Aim:** learn an abstract representation of the environment
 - Use with task-level planners
 - Based on agent behaviours (skills / options)
 - **General:** don't need to be relearned for every new task

Steven James (in collaboration with George Konidaris)

S. James, B. Rosman, G. Konidaris. Learning to Plan with Portable Symbols. ICML/IJCAI/AAMAS 2018 Workshop on Planning and Learning, July 2018.

S. James, B. Rosman, G. Konidaris. Learning Portable Abstract Representations for High-Level Planning. *Under review.*

Requirements: planning with skills

- Learn the preconditions

- Classification problem:

- $P(\text{can execute skill?} \mid \text{current_state})$

- Learn the effects

- Density estimation:

- $P(\text{next_state} \mid \text{current_state}, \text{skill})$

- Possible if options are subgoal i.e.


- $P(\text{next_state} \mid \text{current_state}, \text{skill})$



“SYMBOLS”


Subgoal options

- $P(\text{next_state} \mid \text{current_state}, \text{skill})$
- Partition skills to ensure property holds
 - e.g. “walk to nearest door”



Generating symbols from skills


[Konidaris, 2018]



- Results in abstract MDP/propositional PPDDL
- But $P(s \in I_o)$ and $P(s' | o)$ are distributions/symbols over state space *particular to current task*
 - e.g. grounded in a specific set of xy-coordinates


Towards portability

- Need a representation that facilitates transfer
- Assume agent has sensors which provide it with (lossy) observations
- Augment the state space with action-centric observations
 - **Agent space**
- e.g. robot navigating a building
 - State space: xy-coordinates
 - Agent space: video camera



Portable symbols

- Learning symbols in **agent space**
 - Portable!
 - But: non-Markov and **insufficient for planning**
- Add the subgoal **partition labels** to rules
 - General abstract symbols + grounding → portable rules




Grounding symbols

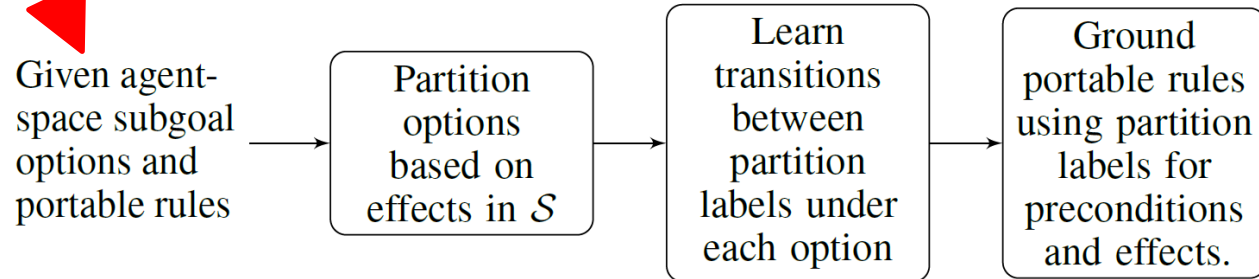
- Learn abstract symbols
- Learning linking functions:
 - Mapping partition numbers from options to their effects
- This gives us a factored MDP or a PPDDL representation
- Provably sufficient for planning

Learning grounded symbols

USING AGENT-SPACE DATA



USING STATE-SPACE DATA



The treasure game



Agent and problem space

- State space: xy -position of agent, key and treasure, angle of levers and state of lock
- Agent space: 9 adjacent cells about the agent




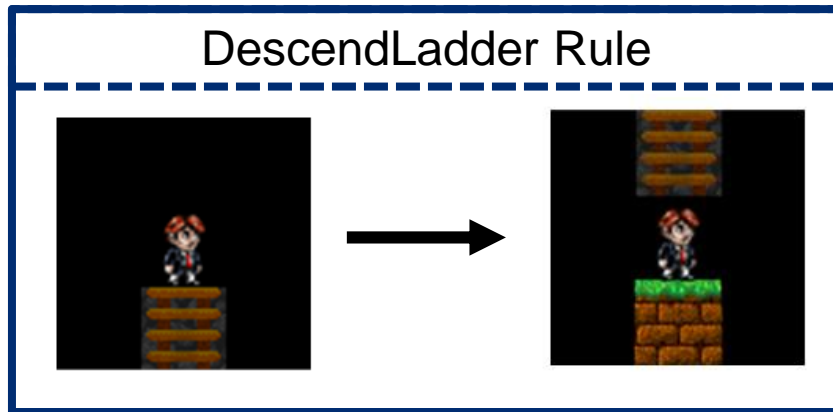
Skills

- Options:
 - GoLeft, GoRight
 - JumpLeft, JumpRight
 - DownRight, DownLeft
 - Interact
 - ClimbLadder,
DescendLadder





Learning portable rules


- Cluster to create subgoal agent-space options
- Use SVM and KDE to estimate preconditions and effects
- Learned rules can be transferred between tasks



Grounding rules

- Partition options in state space to get partition numbers
 - Learn grounded rule instances: linking






Partitioned rules

Precondition:




Negative effect:

Positive effect:

Interact1:




Interact3:



Experiments

- Require fewer samples in subsequent tasks




Portable rules

- Learn abstract rules and their groundings
 - Transfer between domain instances
 - Just by learning linking functions
- But what if there is additional structure?
- In particular, there are many rule instances (objects of interest)?


Ofir Marom

Ofir Marom and Benjamin Rosman. Zero-Shot Transfer with Deictic Object-Oriented Representation in Reinforcement Learning. NIPS, 2018.


Example: Sokoban




Sokoban (legal move)




Sokoban (legal move)



Sokoban (illegal move)




Sokoban (goal)



Representations

$$s = (agent_x = 3, agent_y = 4, box1_x = 4, box1_y = 4, box2_x = 3, box2_y = 2)$$

- Poor scalability
 - 100s of boxes?
 - Transferability?







$$\binom{130}{21} \approx 10^{23}$$

- Effects of actions depend on interactions further away, complicating a mapping to agent space

Object-oriented representations

- Consider objects explicitly
 - Object classes have attributes
 - Relationships based on formal logic: $On(box_1, storage_1)$

Object Classes				
Attributes	(X,Y)	(X,Y)	(X,Y)	(X,Y)
Object Instances	box1(1,5) box2(2,4)


Propositional OO-MDPs_[Duik, 2010]

- Describe transition rules using schemas
- Propositional Object-Oriented MDPs
 - Provably efficient to learn (KWIK bounds)

$$\begin{aligned} & East \wedge \\ & Touch_{East}(Person, Wall) \\ \Rightarrow & Person.x \leftarrow Person.x + 0 \end{aligned}$$

Benefits


- Propositional OO-MDPs
 - Compact representation
 - Efficient learning of rules



Limitations

- Propositional OO-MDPs are efficient, but restrictive


$$\begin{aligned} & East \wedge \\ & Touch_{West}(Box, Person) \wedge Touch_{East}(Box, Wall) \\ & \Rightarrow Box.x \leftarrow ? \end{aligned}$$



Limitations

- Propositional OO-MDPs are efficient, but restrictive
 - Restriction that **preconditions are propositional**
 - Can't refer to the same box

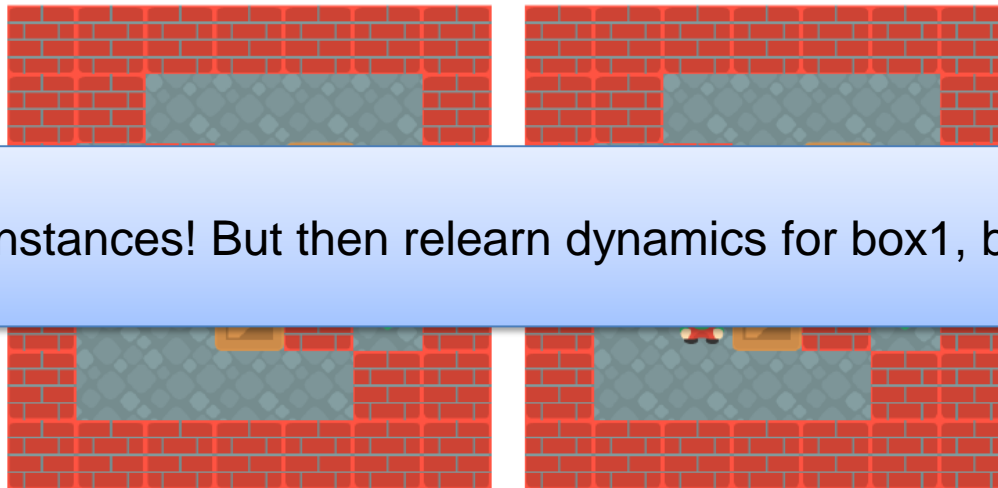
$$\begin{aligned} & \textit{East} \wedge \\ & \textit{Touch}_{\textit{West}}(\textit{Box}, \textit{Person}) \wedge \textit{Touch}_{\textit{East}}(\textit{Box}, \textit{Wall}) \\ & \Rightarrow \textit{Box}.x \leftarrow ? \end{aligned}$$



Limitations

- Propositional OO-MDPs are efficient, but restrictive
 - Restriction that **preconditions are propositional**
 - Can't refer to the same box

$$\begin{aligned} & East \wedge \\ & Touch_{West}(Box, Person) \wedge Touch_{East}(Box, Wall) \\ & \Rightarrow Box.x \leftarrow ? \end{aligned}$$



Ground instances! But then relearn dynamics for box1, box2, etc.

Deictic OO-MDPs

- **Deictic predicates** instead of propositions
 - Grounded only with respect to a central deictic object (“*me*” or “*this*”)
 - Relates to other non-grounded objects
- Transition dynamics of $Box.x$ depends on grounded box object


$$\begin{aligned} & East \wedge \\ & Touch_{West}(box, Person) \wedge Touch_{East}(box, Wall) \\ & \Rightarrow box.x \leftarrow box.x + 0 \end{aligned}$$

- Also provably efficient


Learning the dynamics

- Learning from experience:
 - For each action, how do attributes change?
- KWIK framework
- Propositional OO-MDPs: DOORMAX algorithm
 - Transition dynamics for each attribute and action must be representable as a binary tree
 - Effects at the leaf nodes
 - Each possible effect can occur at most at one leaf, except for a failure condition (globally nothing changes)

Learning the dynamics





Example: action = North




- Given:
 - Each effect can only occur once on the tree
 - Global failure condition
 - Deterministic effects
- Learn from common elements in state propositions (experience)

DOORMAX for deictic OO-MDPs



- We adapt the DOORMAX algorithm to deictic OO-MDPs
 - Remove global failure condition
 - Bound the number of times a condition can occur
 - Can still be learned efficiently



Experiments



$\sim 8k$ states



$\sim 1M$ states

- Zero-shot transfer: one run of value iteration

Experiments


- Taxi domain
 - Multi-passengers
 - Only one in the taxi at a time
- On executing a *pickup* action
 - Change *in_taxi* attribute of **correct** passenger



Experiments



1 passenger



2 passengers



3 passengers



4 passengers

Take away thoughts

- **Reinforcement learning** gives us a powerful tool for learning behaviours, but extra work is required for generalisation
- Reasoning in an **agent-centric manner**:
 - Symbol-based view on skills
 - Enable knowledge reuse
- Reasoning in an **object-centric manner**:
 - Learn models of local object interactions
 - Efficient learning and transfer

Thank you!



www.benjaminrosman.com – www.raillab.org